


Make every AI token count.

TrimToken AI is the intelligent gateway that routes each request to the best-fit model for quality, latency, and cost.

**A practical response to one question:
Why pay frontier-model prices for every prompt?**



The AI cost problem is becoming a systems problem.

Teams are adding AI faster than they can control spend.

01 One-size-fits-all model selection

Simple requests are often sent to expensive frontier models.

02 Prices and capabilities keep moving

Provider pricing, latency, and quality change constantly.

03 Developers need a drop-in path

Cost optimization cannot require rewriting every AI product.

OUR INSIGHT

**The best model is not always the biggest model.
It is the model that meets the task's quality bar at the lowest total cost**

TrimToken turns model choice into an always-on optimization layer.

TrimToken AI: an intelligent control plane for LLM traffic.

A gateway that classifies each query, compares live model economics, and dispatches to the best-fit provider.

?

Query intelligence

Fine-tuned Qwen routing classifier scores complexity in sub-millisecond time.

?

Price-aware routing

Pareto utility balances quality, latency, and cost in real time.

?

Production resilience

OpenAI-compatible API, failover mesh, and observability built in.

Designed for enterprise AI teams that need lower spend without sacrificing the user experience.

From prompt to outcome in four routing decisions.

1

Ingest

Standard OpenAI-compatible request enters the gateway.

2

Classify

Query is labeled by complexity and task type.

3

Optimize

Live pricing and policy constraints rank candidates.

4

Execute

Best-fit model responds; failover protects reliability.

ROUTING POLICY EXAMPLE

**If a request is simple,
choose the lowest-cost model
that clears the quality
threshold.**

Deep reasoning and code move up the model
spectrum only when needed.

QUALITY BAR + COST CEILING

Why customers pay: measurable savings, not another model.

TrimToken sits between the application and model providers, creating a measurable cost and reliability layer.

UP TO 70%

targeted reduction in
enterprise LLM API overhead

< 1.2 ms

routing-classifier latency
target

1-line swap

OpenAI-compatible gateway
integration

The economic moat is the routing data: every query improves policy, benchmarks, and model economics.

A recurring SaaS model aligned with customer savings.

Gateway

Usage-based

Per-request fee or percentage of savings for teams getting started.

Scale

Platform subscription

Monthly plan for policies, analytics, budgets, and team controls.

Enterprise

Annual contract

Private deployment, SLAs, custom policies, security reviews, and support.

The more traffic a customer routes, the more value we create ? and the more recurring revenue we earn.

Expansion: observability ? governance ? private routing models ? multi-cloud optimization

Genesis can help us turn a working prototype into a category platform.

We are seeking Genesis support to validate the product with design partners and build the data moat around cost-aware routing.


- 40%** **Pilot deployments**
Design partners and production integrations.
- 30%** **Routing intelligence**
Benchmarks, policies, and fine-tuned models.
- 20%** **Reliability + security**
Observability, failover, and enterprise controls.
- 10%** **GTM**
ROI tooling, case studies, and onboarding.

THE MILESTONE

Prove repeatable savings in production.

A focused path from prototype to paid, measurable deployments.

The future: a routing intelligence layer for every AI stack.



Thank you ? TrimToken AI ? Let's make AI economics sustainable.